

Feature Selection

2022-04-01

```
# Load Libraries  
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.1.3
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v ggplot2 3.3.5      v purrr  0.3.4  
## v tibble  3.1.6      v dplyr  1.0.8  
## v tidyr   1.2.0      v stringr 1.4.0  
## v readr   2.1.2      v forcats 0.5.1
```

```
## Warning: package 'ggplot2' was built under R version 4.1.3
```

```
## Warning: package 'dplyr' was built under R version 4.1.3
```

```
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()
```

```
library(moments)  
library(heatmaply)
```

```
## Warning: package 'heatmaply' was built under R version 4.1.3
```

```
## Loading required package: plotly
```

```
##  
## Attaching package: 'plotly'
```

```
## The following object is masked from 'package:ggplot2':  
##  
## last_plot
```

```
## The following object is masked from 'package:stats':  
##  
## filter
```

```
## The following object is masked from 'package:graphics':  
##  
## layout
```

```
## Loading required package: viridis

## Warning: package 'viridis' was built under R version 4.1.3

## Loading required package: viridisLite

##
## =====
## Welcome to heatmaply version 1.3.0
##
## Type citation('heatmaply') for how to cite the package.
## Type ?heatmaply for the main documentation.
##
## The github page is: https://github.com/talgilili/heatmaply/
## Please submit your suggestions and bug-reports at: https://github.com/talgilili/heatmaply/issues
## You may ask questions at stackoverflow, use the r and heatmaply tags:
##   https://stackoverflow.com/questions/tagged/heatmaply
## =====
```

```
library(dummies)
```

```
## dummies-1.5.6 provided by Decision Patterns
```

```
# Import the dataset
smarket <- read.csv("http://bit.ly/CarreFourDataset")
```

```
# Preview the first few rows
head(smarket)
```

```
##   Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 1 750-67-8428      A      Member Female   Health and beauty      74.69
## 2 226-31-3081      C      Normal Female Electronic accessories    15.28
## 3 631-41-3108      A      Normal  Male    Home and lifestyle     46.33
## 4 123-19-1176      A      Member  Male    Health and beauty     58.22
## 5 373-73-7910      A      Normal  Male    Sports and travel      86.31
## 6 699-14-3026      C      Normal  Male    Electronic accessories   85.39
##   Quantity      Tax      Date Time      Payment      cogs gross.margin.percentage
## 1         7 26.1415 1/5/2019 13:08      Ewallet 522.83          4.761905
## 2         5  3.8200 3/8/2019 10:29        Cash  76.40          4.761905
## 3         7 16.2155 3/3/2019 13:23 Credit card 324.31          4.761905
## 4         8 23.2880 1/27/2019 20:33      Ewallet 465.76          4.761905
## 5         7 30.2085 2/8/2019 10:37      Ewallet 604.17          4.761905
## 6         7 29.8865 3/25/2019 18:30      Ewallet 597.73          4.761905
##   gross.income Rating      Total
## 1      26.1415      9.1 548.9715
## 2       3.8200      9.6  80.2200
## 3      16.2155      7.4 340.5255
## 4      23.2880      8.4 489.0480
## 5      30.2085      5.3 634.3785
## 6      29.8865      4.1 627.6165
```

```
# Preview the last few rows
tail(smarket)
```

```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 995  652-49-6720      C      Member Female Electronic accessories    60.95
## 996  233-67-5758      C      Normal  Male   Health and beauty      40.35
## 997  303-96-2227      B      Normal Female Home and lifestyle    97.38
## 998  727-02-1313      A      Member  Male   Food and beverages     31.84
## 999  347-56-2442      A      Normal  Male   Home and lifestyle     65.82
## 1000 849-09-3807      A      Member Female Fashion accessories    88.34
##      Quantity      Tax      Date Time Payment  cogs gross.margin.percentage
## 995          1  3.0475 2/18/2019 11:40 Ewallet  60.95          4.761905
## 996          1  2.0175 1/29/2019 13:46 Ewallet  40.35          4.761905
## 997         10 48.6900 3/2/2019 17:16 Ewallet 973.80          4.761905
## 998          1  1.5920 2/9/2019 13:22   Cash  31.84          4.761905
## 999          1  3.2910 2/22/2019 15:33   Cash  65.82          4.761905
## 1000         7 30.9190 2/18/2019 13:28   Cash 618.38          4.761905
##      gross.income Rating      Total
## 995          3.0475      5.9    63.9975
## 996          2.0175      6.2    42.3675
## 997         48.6900      4.4  1022.4900
## 998          1.5920      7.7    33.4320
## 999          3.2910      4.1    69.1110
## 1000         30.9190      6.6   649.2990
```

```
# Check number of records and variables
dim(smarket)
```

```
## [1] 1000  16
```

We have 1,000 records and 16 variables

```
# Check the datatypes of our dataset
glimpse(smarket)
```

```
## Rows: 1,000
## Columns: 16
## $ Invoice.ID      <chr> "750-67-8428", "226-31-3081", "631-41-3108", "~
## $ Branch         <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A", "~
## $ Customer.type  <chr> "Member", "Normal", "Normal", "Member", "Norma~
## $ Gender         <chr> "Female", "Female", "Male", "Male", "Male", "M~
## $ Product.line   <chr> "Health and beauty", "Electronic accessories",~
## $ Unit.price     <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 68.8~
## $ Quantity       <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10, 10~
## $ Tax            <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085, 29~
## $ Date           <chr> "1/5/2019", "3/8/2019", "3/3/2019", "1/27/2019~
## $ Time           <chr> "13:08", "10:29", "13:23", "20:33", "10:37", "~
## $ Payment        <chr> "Ewallet", "Cash", "Credit card", "Ewallet", "~
## $ cogs           <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597.73,~
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7619~
## $ gross.income   <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085, 29~
## $ Rating         <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2, 5~
## $ Total          <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634.378~
```

```
# Check the summary of our dataset
summary(smarket)
```

```
## Invoice.ID           Branch           Customer.type           Gender
## Length:1000         Length:1000         Length:1000         Length:1000
## Class :character     Class :character     Class :character     Class :character
## Mode :character      Mode :character      Mode :character      Mode :character
##
##
##
## Product.line         Unit.price          Quantity            Tax
## Length:1000         Min. :10.08         Min. : 1.00         Min. : 0.5085
## Class :character     1st Qu.:32.88       1st Qu.: 3.00       1st Qu.: 5.9249
## Mode :character      Median :55.23        Median : 5.00       Median :12.0880
##                      Mean :55.67          Mean : 5.51         Mean :15.3794
##                      3rd Qu.:77.94       3rd Qu.: 8.00       3rd Qu.:22.4453
##                      Max. :99.96          Max. :10.00         Max. :49.6500
##
## Date                 Time                 Payment             cogs
## Length:1000         Length:1000         Length:1000         Min. : 10.17
## Class :character     Class :character     Class :character     1st Qu.:118.50
## Mode :character      Mode :character     Mode :character     Median :241.76
##                      Mean :307.59
##                      3rd Qu.:448.90
##                      Max. :993.00
##
## gross.margin.percentage gross.income          Rating              Total
## Min. :4.762          Min. : 0.5085         Min. : 4.000         Min. : 10.68
## 1st Qu.:4.762        1st Qu.: 5.9249       1st Qu.: 5.500       1st Qu.:124.42
## Median :4.762        Median :12.0880       Median : 7.000       Median :253.85
## Mean :4.762          Mean :15.3794         Mean : 6.973         Mean :322.97
## 3rd Qu.:4.762        3rd Qu.:22.4453       3rd Qu.: 8.500       3rd Qu.:471.35
## Max. :4.762          Max. :49.6500         Max. :10.000        Max. :1042.65
```

```
# Check the column names
names(smarket)
```

```
## [1] "Invoice.ID"           "Branch"
## [3] "Customer.type"        "Gender"
## [5] "Product.line"         "Unit.price"
## [7] "Quantity"             "Tax"
## [9] "Date"                 "Time"
## [11] "Payment"              "cogs"
## [13] "gross.margin.percentage" "gross.income"
## [15] "Rating"               "Total"
```

Data Cleaning

```
# Let's check for missing values
colSums(is.na(smarket))
```

```
## Invoice.ID           Branch           Customer.type
```

```
##          0          0          0
##      Gender      Product.line      Unit.price
##          0          0          0
##      Quantity      Tax      Date
##          0          0          0
##      Time      Payment      cogs
##          0          0          0
## gross.margin.percentage      gross.income      Rating
##          0          0          0
##      Total
##          0
```


We have no missing values.

```
# Check for duplicate values
sum(duplicated(smarket))
```


```
## [1] 0
```

We have no duplicate values


```
# Checking for outliers in our numerical variables
boxplot(smarket[, c(6,7)])
```




```
# Checking for outliers in our numerical variables  
boxplot(smarket[, c(8,12)])
```



```
# Checking for outliers in our numerical variables  
boxplot(smarket[, c(13,14)])
```



```
# Checking for outliers in our numerical variables  
boxplot(smarket[, c(15,16)])
```



There is presence of outliers in Tax, Cogs, Gross income and Total Variables

```
# Convert the date column from character datatype to date datatype
smarket$Date <- as.Date(smarket$Date, "%m/%d/%Y")
# Confirm changes made
glimpse(smarket)
```

```
## Rows: 1,000
## Columns: 16
## $ Invoice.ID      <chr> "750-67-8428", "226-31-3081", "631-41-3108", "~
## $ Branch         <chr> "A", "C", "A", "A", "A", "C", "A", "C", "A", "~
## $ Customer.type  <chr> "Member", "Normal", "Normal", "Member", "Norma~
## $ Gender         <chr> "Female", "Female", "Male", "Male", "Male", "M~
## $ Product.line   <chr> "Health and beauty", "Electronic accessories",~
## $ Unit.price     <dbl> 74.69, 15.28, 46.33, 58.22, 86.31, 85.39, 68.8~
## $ Quantity       <int> 7, 5, 7, 8, 7, 7, 6, 10, 2, 3, 4, 4, 5, 10, 10~
## $ Tax            <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085, 29~
## $ Date           <date> 2019-01-05, 2019-03-08, 2019-03-03, 2019-01-2~
## $ Time           <chr> "13:08", "10:29", "13:23", "20:33", "10:37", "~
## $ Payment        <chr> "Ewallet", "Cash", "Credit card", "Ewallet", "~
## $ cogs           <dbl> 522.83, 76.40, 324.31, 465.76, 604.17, 597.73,~
## $ gross.margin.percentage <dbl> 4.761905, 4.761905, 4.761905, 4.761905, 4.7619~
## $ gross.income   <dbl> 26.1415, 3.8200, 16.2155, 23.2880, 30.2085, 29~
## $ Rating         <dbl> 9.1, 9.6, 7.4, 8.4, 5.3, 4.1, 5.8, 8.0, 7.2, 5~
## $ Total          <dbl> 548.9715, 80.2200, 340.5255, 489.0480, 634.378~
```



```
# We will extract month from the date column
smarket$Month<- format(smarket$Date, "%m")
head(smarket)
```

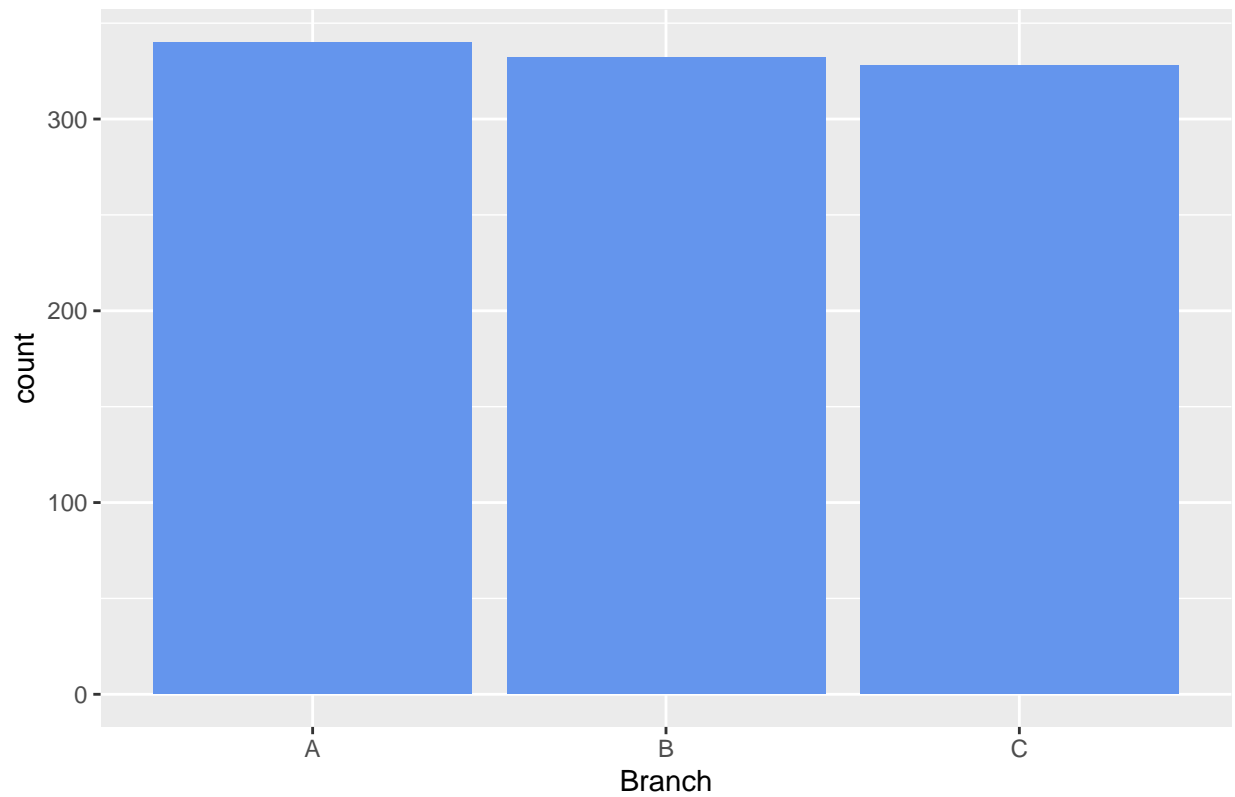
```
##      Invoice.ID Branch Customer.type Gender      Product.line Unit.price
## 1 750-67-8428      A      Member Female      Health and beauty      74.69
## 2 226-31-3081      C      Normal Female Electronic accessories      15.28
## 3 631-41-3108      A      Normal  Male      Home and lifestyle      46.33
## 4 123-19-1176      A      Member  Male      Health and beauty      58.22
## 5 373-73-7910      A      Normal  Male      Sports and travel      86.31
## 6 699-14-3026      C      Normal  Male Electronic accessories      85.39
##      Quantity      Tax      Date Time      Payment      cogs gross.margin.percentage
## 1          7 26.1415 2019-01-05 13:08      Ewallet 522.83          4.761905
## 2          5  3.8200 2019-03-08 10:29      Cash 76.40          4.761905
## 3          7 16.2155 2019-03-03 13:23 Credit card 324.31          4.761905
## 4          8 23.2880 2019-01-27 20:33      Ewallet 465.76          4.761905
## 5          7 30.2085 2019-02-08 10:37      Ewallet 604.17          4.761905
## 6          7 29.8865 2019-03-25 18:30      Ewallet 597.73          4.761905
##      gross.income Rating      Total Month
## 1      26.1415      9.1 548.9715      01
## 2       3.8200      9.6  80.2200      03
## 3      16.2155      7.4 340.5255      03
## 4      23.2880      8.4 489.0480      01
## 5      30.2085      5.3 634.3785      02
## 6      29.8865      4.1 627.6165      03
```

```
# Replace the numbers in months with the name of the month
smarket$Month[smarket$Month == "01"] <- "January"
smarket$Month[smarket$Month == "02"] <- "February"
smarket$Month[smarket$Month == "03"] <- "March"
```

Univariate Analysis


```
# Bar chart on revenue
gg.1 <- ggplot (data = smarket, aes(x= Branch)) +
  geom_bar(fill = "cornflowerblue")
gg.1 + ggtitle("Bar chart on Branch")
```

Bar chart on Branch




The number of branches in our data are not too different from each other

```
# Bar Chart on Customer Type
gg.2 <- ggplot (data = smarket, aes(x= Customer.type)) +
  geom_bar(fill = "cornflowerblue")
gg.2 + ggtitle("Bar chart on Customer Type")
```



There isn't a big difference in the number of member and normal customers

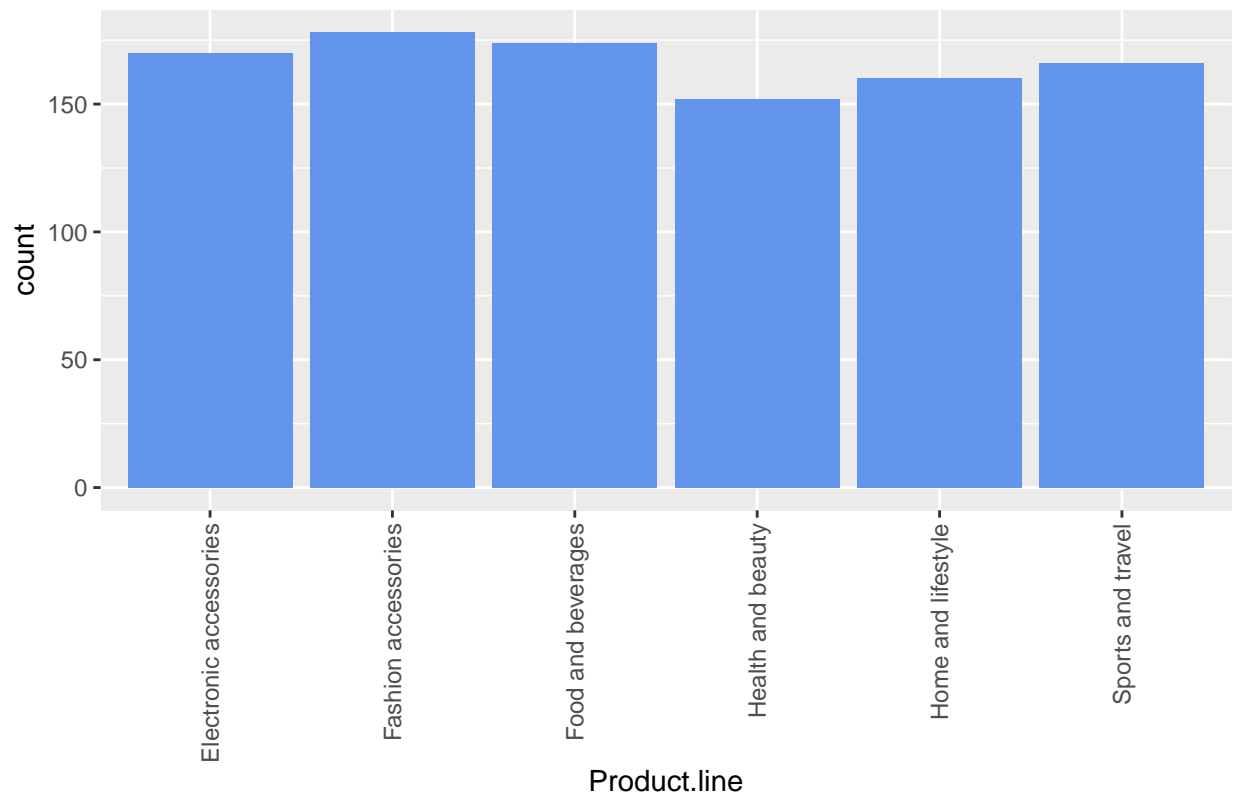
```
# Bar Chart on Gender  
gg.3 <- ggplot (data = smarket, aes(x= Gender)) +  
  geom_bar(fill = "cornflowerblue")  
gg.3 + ggtitle("Bar chart on Gender")
```



There isn't a big difference in the number of male and female customers

```
# Bar Chart on Product Line
gg.4 <- ggplot (data = smarket, aes(x= Product.line)) +
  geom_bar(fill = "cornflowerblue") + theme(axis.text.x = element_text(
    angle = 90, vjust = .5, hjust = 1
  ))
gg.4 + ggtitle("Bar chart on Product Line")
```

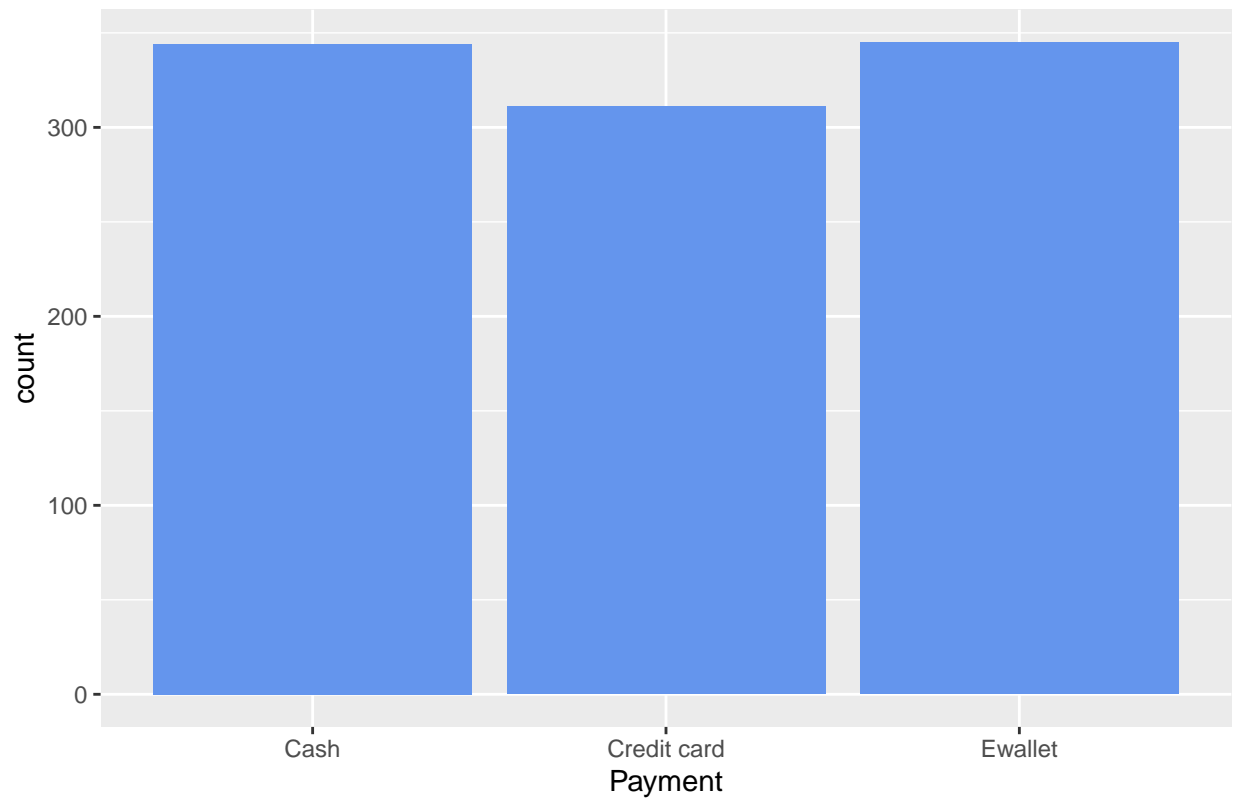
Bar chart on Product Line



Fashion accessories is the most popular product followed closely by food and beverages

```
# Bar Chart on Payment
gg.5 <- ggplot (data = smarket, aes(x= Payment)) +
  geom_bar(fill = "cornflowerblue")
gg.5 + ggtitle("Bar chart on Payment")
```


Bar chart on Payment



There is no big difference in the number of customers who paid via Ewallet and cash

```
# Bar Chart on Month  
gg.6<- ggplot (data = smarket, aes(x= Month)) +  
  geom_bar(fill = "cornflowerblue")  
gg.6 + ggtitle("Bar chart on Month")
```

Bar chart on Month



January had the most transactions followed closely by March

```
# Frequency chart on Unit price
freq.1 <- table(smarket$Unit.price)
head(sort(freq.1, decreasing = T))
```

```
##
## 83.77  15.5  15.8 18.08 19.15 20.01
##      3    2    2    2    2    2
```

```
# Frequency chart on Quantity
freq.2 <- table(smarket$Quantity)
head(sort(freq.2, decreasing = T))
```

```
##
## 10  1  4  5  7  6
## 119 112 109 102 102 98
```

```
# Frequency chart on Tax
freq.3 <- table(smarket$Tax)
head(sort(freq.3, decreasing = T), n=15)
```

```
##
##  4.154  4.464  8.377  9.0045 10.326 10.3635 12.57 13.188 22.428 39.48
##      2      2      2      2      2      2      2      2      2      2
```

```
## 0.5085 0.6045 0.627 0.639 0.699
##      1      1      1      1      1
```

```
# Frequency chart on cogs
freq.4 <- table(smarket$cogs)
head(sort(freq.4, decreasing = T), n=15)
```

```
##
## 83.08 89.28 167.54 180.09 206.52 207.27 251.4 263.76 448.56 789.6 10.17
##      2      2      2      2      2      2      2      2      2      2      1
## 12.09 12.54 12.78 13.98
##      1      1      1      1
```

```
# Frequency chart on gross margin percentage
freq.5 <- table(smarket$gross.margin.percentage)
head(sort(freq.5, decreasing = T))
```

```
## 4.761904762
##      1000
```

```
# Frequency chart on Rating
freq.5 <- table(smarket$Rating)
head(sort(freq.5, decreasing = T))
```

```
##
## 6 6.6 4.2 9.5 5 5.1
## 26 24 22 22 21 21
```

```
# Frequency chart on Gross Income
freq.5 <- table(smarket$gross.income)
head(sort(freq.5, decreasing = T), n=10)
```

```
##
## 4.154 4.464 8.377 9.0045 10.326 10.3635 12.57 13.188 22.428 39.48
##      2      2      2      2      2      2      2      2      2      2
```

```
# Frequency chart on Total
freq.5 <- table(smarket$Total)
head(sort(freq.5, decreasing = T), n=15)
```

```
##
## 87.234 93.744 175.917 189.0945 216.846 217.6335 263.97 276.948
##      2      2      2      2      2      2      2      2
## 470.988 829.08 10.6785 12.6945 13.167 13.419 14.679
##      2      2      1      1      1      1      1
```

```
# Identify the mean, median, min, max and quantile of our numerical variables
summary(smarket[,c(6:8,12:16)])
```



```
##      Unit.price      Quantity      Tax      cogs
##  Min.   :10.08   Min.   : 1.00   Min.   : 0.5085   Min.   : 10.17
##  1st Qu.:32.88   1st Qu.: 3.00   1st Qu.: 5.9249   1st Qu.:118.50
##  Median :55.23   Median : 5.00   Median :12.0880   Median :241.76
##  Mean   :55.67   Mean   : 5.51   Mean   :15.3794   Mean   :307.59
##  3rd Qu.:77.94   3rd Qu.: 8.00   3rd Qu.:22.4453   3rd Qu.:448.90
##  Max.   :99.96   Max.   :10.00   Max.   :49.6500   Max.   :993.00
##  gross.margin.percentage gross.income      Rating      Total
##  Min.   :4.762      Min.   : 0.5085   Min.   : 4.000   Min.   : 10.68
##  1st Qu.:4.762      1st Qu.: 5.9249   1st Qu.: 5.500   1st Qu.: 124.42
##  Median :4.762      Median :12.0880   Median : 7.000   Median : 253.85
##  Mean   :4.762      Mean   :15.3794   Mean   : 6.973   Mean   : 322.97
##  3rd Qu.:4.762      3rd Qu.:22.4453   3rd Qu.: 8.500   3rd Qu.: 471.35
##  Max.   :4.762      Max.   :49.6500   Max.   :10.000   Max.   :1042.65
```

```
# We'll also check the mode of our numerical variables
```

```
mode <- function(x) {
  u <- unique(x)
  tab <- tabulate(match(x, u))
  u[tab == max(tab)]
}
```

```
print("Mode of unit price is:")
```

```
## [1] "Mode of unit price is:"
```

```
mode(smarket$Unit.price)
```

```
## [1] 83.77
```

```
print("Mode of tax is:")
```

```
## [1] "Mode of tax is:"
```

```
mode(smarket$Tax)
```

```
## [1] 39.4800  9.0045 10.3260 12.5700 10.3635 13.1880  4.1540  8.3770 22.4280
## [10]  4.4640
```

```
print("Mode of COGS is:")
```

```
## [1] "Mode of COGS is:"
```

```
mode(smarket$cogs)
```

```
## [1] 789.60 180.09 206.52 251.40 207.27 263.76  83.08 167.54 448.56  89.28
```

```
print("Mode of gross margin percentage is:")
```

```
## [1] "Mode of gross margin percentage is:"
```

```
mode(smarket$gross.margin.percentage)
```

```
## [1] 4.761905
```

```
print("Mode of Gross income:")
```

```
## [1] "Mode of Gross income:"
```

```
mode(smarket$gross.income)
```

```
## [1] 39.4800 9.0045 10.3260 12.5700 10.3635 13.1880 4.1540 8.3770 22.4280
```

```
## [10] 4.4640
```

```
print("Mode of Rating is:")
```

```
## [1] "Mode of Rating is:"
```

```
mode(smarket$Rating)
```

```
## [1] 6
```

```
print("Mode of Total is:")
```

```
## [1] "Mode of Total is:"
```

```
mode(smarket$Total)
```

```
## [1] 829.0800 189.0945 216.8460 263.9700 217.6335 276.9480 87.2340 175.9170
```

```
## [9] 470.9880 93.7440
```

```
# Check the variance of our numerical variables
```

```
print("Variance of unit price is:")
```

```
## [1] "Variance of unit price is:"
```

```
var(smarket$Unit.price)
```

```
## [1] 701.9653
```

```
print("Variance of tax is:")
```

```
## [1] "Variance of tax is:"
```

```
var(smarket$Tax)
```

```
## [1] 137.0966
```

```
print("Variance of COGS is:")
```

```
## [1] "Variance of COGS is:"
```

```
var(smarket$cogs)
```

```
## [1] 54838.64
```

```
print("Variance of gross margin percentage is:")
```

```
## [1] "Variance of gross margin percentage is:"
```

```
var(smarket$gross.margin.percentage)
```

```
## [1] 0
```

```
print("Variance of Gross income:")
```

```
## [1] "Variance of Gross income:"
```

```
var(smarket$gross.income)
```

```
## [1] 137.0966
```

```
print("Variance of Rating is:")
```

```
## [1] "Variance of Rating is:"
```

```
var(smarket$Rating)
```

```
## [1] 2.953518
```

```
print("Variance of Total is:")
```

```
## [1] "Variance of Total is:"
```

```
var(smarket$Total)
```

```
## [1] 60459.6
```

```
# Check the skewness of our variables  
print("The skewness of our variables:")
```

```
## [1] "The skewness of our variables:"
```

```
sapply(smarket[,c(6:8,12:16)],skewness)
```

```
##           Unit.price           Quantity           Tax  
##           0.007066827           0.012921628           0.891230392  
##           cogs gross.margin.percentage           gross.income  
##           0.891230392           NaN           0.891230392  
##           Rating           Total  
##           0.008996129           0.891230392
```

```
# Check the kurtosis of our variables  
print("The kurtosis of our variables:")
```





```
## [1] "The kurtosis of our variables:"
```

```
sapply(smarket[,c(6:8,12:16)],kurtosis)
```

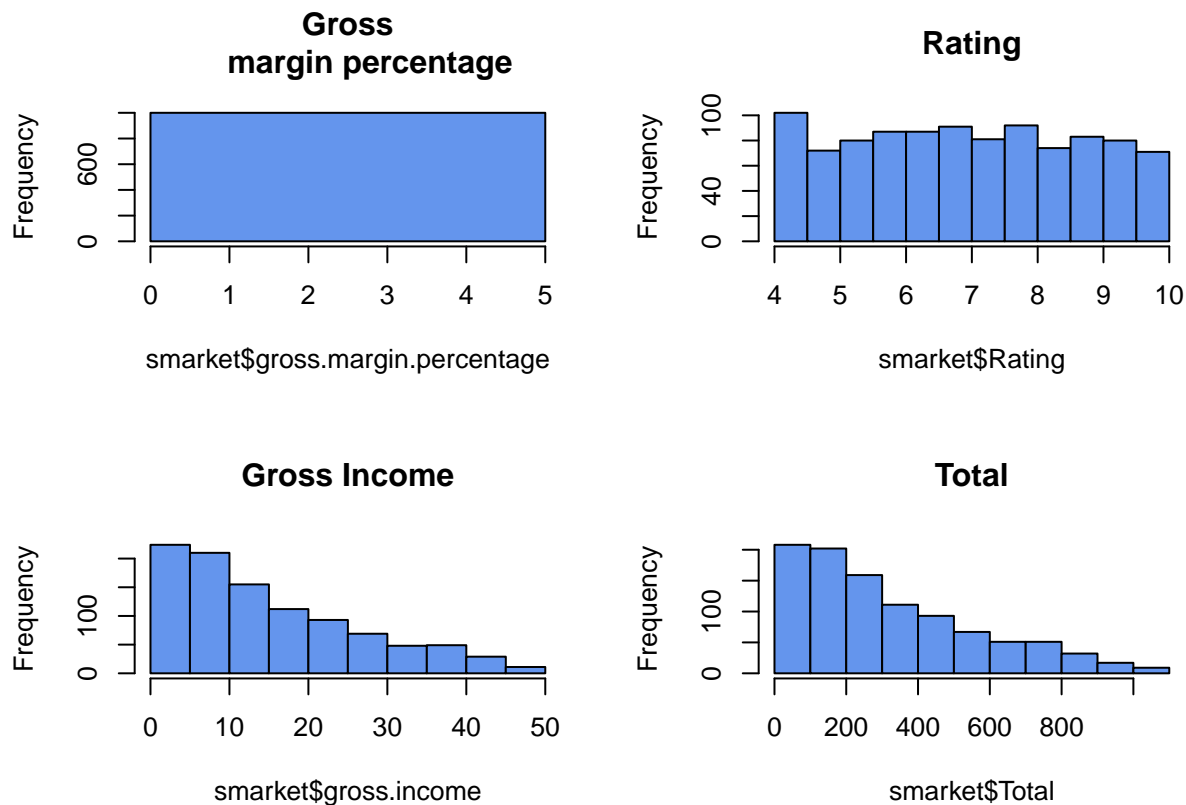
```
##           Unit.price           Quantity           Tax  
##           1.781499           1.784528           2.912530  
##           cogs gross.margin.percentage           gross.income  
##           2.912530           NaN           2.912530  
##           Rating           Total  
##           1.848169           2.912530
```

Our numerical variables do not have heavy tails

```
# Plot histograms of our numeric variables  
par(mfcol=c(2,2))  
hist(smarket$Unit.price, col = "cornflowerblue", main = "Unit Price")  
hist(smarket$Quantity, col = "cornflowerblue",  
     main = "Quantity")  
hist(smarket$Tax, col = "cornflowerblue", main="Tax")  
hist(smarket$cogs, col = "cornflowerblue",  
     main ="COGS")
```



```
# Plot histograms of our numeric variables
par(mfcol=c(2,2))
hist(smarket$gross.margin.percentage, col = "cornflowerblue", main = "Gross
    margin percentage")
hist(smarket$gross.income, col = "cornflowerblue",
    main = "Gross Income")
hist(smarket$Rating, col = "cornflowerblue", main="Rating")
hist(smarket$Total, col = "cornflowerblue",
    main ="Total")
```



Unit price quantity and Rating are symmetrical and the rest are moderately skewed.

Bivariate Analysis

```
# Assess the correlation of our numerical variables
cor(smarket[,c(6:8,12:16)])
```

```
## Warning in cor(smarket[, c(6:8, 12:16)]): the standard deviation is zero
```

```
##           Unit.price  Quantity      Tax      cogs
## Unit.price      1.00000000  0.01077756  0.6339621  0.6339621
## Quantity        0.01077756  1.00000000  0.7055102  0.7055102
## Tax             0.633962089  0.70551019  1.0000000  1.0000000
## cogs            0.633962089  0.70551019  1.0000000  1.0000000
## gross.margin.percentage      NA      NA      NA      NA
## gross.income      0.633962089  0.70551019  1.0000000  1.0000000
## Rating          -0.008777507 -0.01581490 -0.0364417 -0.0364417
## Total           0.633962089  0.70551019  1.0000000  1.0000000
## gross.margin.percentage gross.income      Rating
## Unit.price              NA      0.6339621 -0.008777507
## Quantity                NA      0.7055102 -0.015814905
## Tax                    NA      1.0000000 -0.036441705
## cogs                   NA      1.0000000 -0.036441705
## gross.margin.percentage      1      NA      NA
```

```
## gross.income      NA      1.0000000 -0.036441705
## Rating            NA      -0.0364417  1.000000000
## Total             NA      1.0000000 -0.036441705
##                  Total
## Unit.price        0.6339621
## Quantity          0.7055102
## Tax               1.0000000
## cogs              1.0000000
## gross.margin.percentage NA
## gross.income      1.0000000
## Rating            -0.0364417
## Total             1.0000000
```

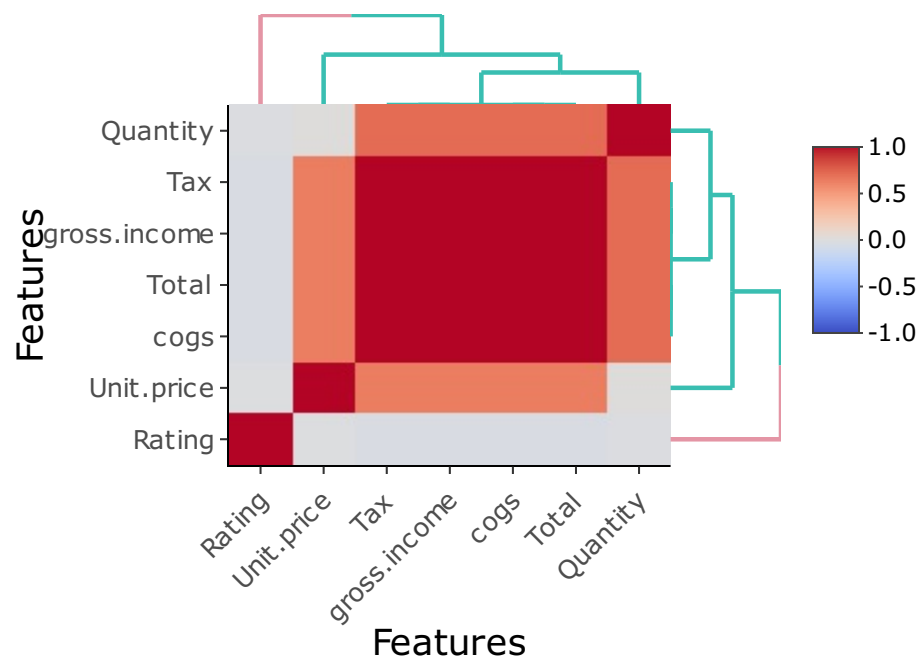
```
# Check for correlation without the gross.margin.percentage
cor(smarket[,c(6:8,12,14,15,16)])
```

```
##          Unit.price  Quantity      Tax      cogs gross.income
## Unit.price  1.000000000  0.01077756  0.6339621  0.6339621  0.6339621
## Quantity    0.010777564  1.00000000  0.7055102  0.7055102  0.7055102
## Tax         0.633962089  0.70551019  1.0000000  1.0000000  1.0000000
## cogs        0.633962089  0.70551019  1.0000000  1.0000000  1.0000000
## gross.income 0.633962089  0.70551019  1.0000000  1.0000000  1.0000000
## Rating      -0.008777507 -0.01581490 -0.0364417 -0.0364417 -0.0364417
## Total       0.633962089  0.70551019  1.0000000  1.0000000  1.0000000
##          Rating      Total
## Unit.price -0.008777507  0.6339621
## Quantity   -0.015814905  0.7055102
## Tax        -0.036441705  1.0000000
## cogs       -0.036441705  1.0000000
## gross.income -0.036441705  1.0000000
## Rating      1.000000000 -0.0364417
## Total      -0.036441705  1.0000000
```

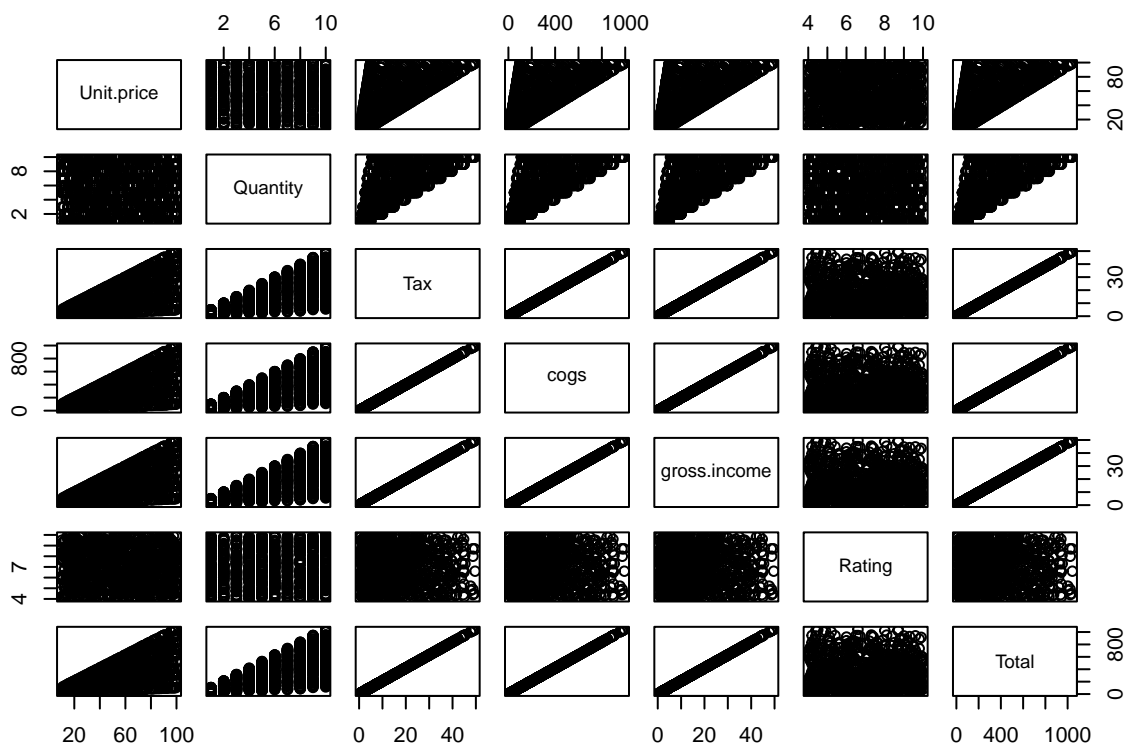
There is a strong positive correlation between quantity and Tax ,cog, Gross income and total Which makes sense because Quantity is used in the calculation of these variables

There's a weak negative correlation between Rating and Tax ,cog, Gross income and total

```
# Plotting corr heatmap
heatmaply_cor(x = cor(smarket[,c(6:8,12,14:16)]), xlab = "Features",
              ylab = "Features", k_col = 2, k_row = 2)
```



```
# Plot a pairplot
pairs(smarket[,c(6:8,12,14:16)])
```





```
# Bar Chart of Branch vs Total
```

```
agg_1 <- aggregate(smarket$Total,list(smarket$Branch),sum)
agg_1
```

```
##   Group.1      x
## 1      A 106200.4
## 2      B 106197.7
## 3      C 110568.7
```

```
ggplot(data = agg_1 , aes(x= Group.1, y =x)) + geom_col(fill ="cornflowerblue") +
  labs(y="Total", x="Branch")+ggtitle("Bar chart of Branch vs Total")
```

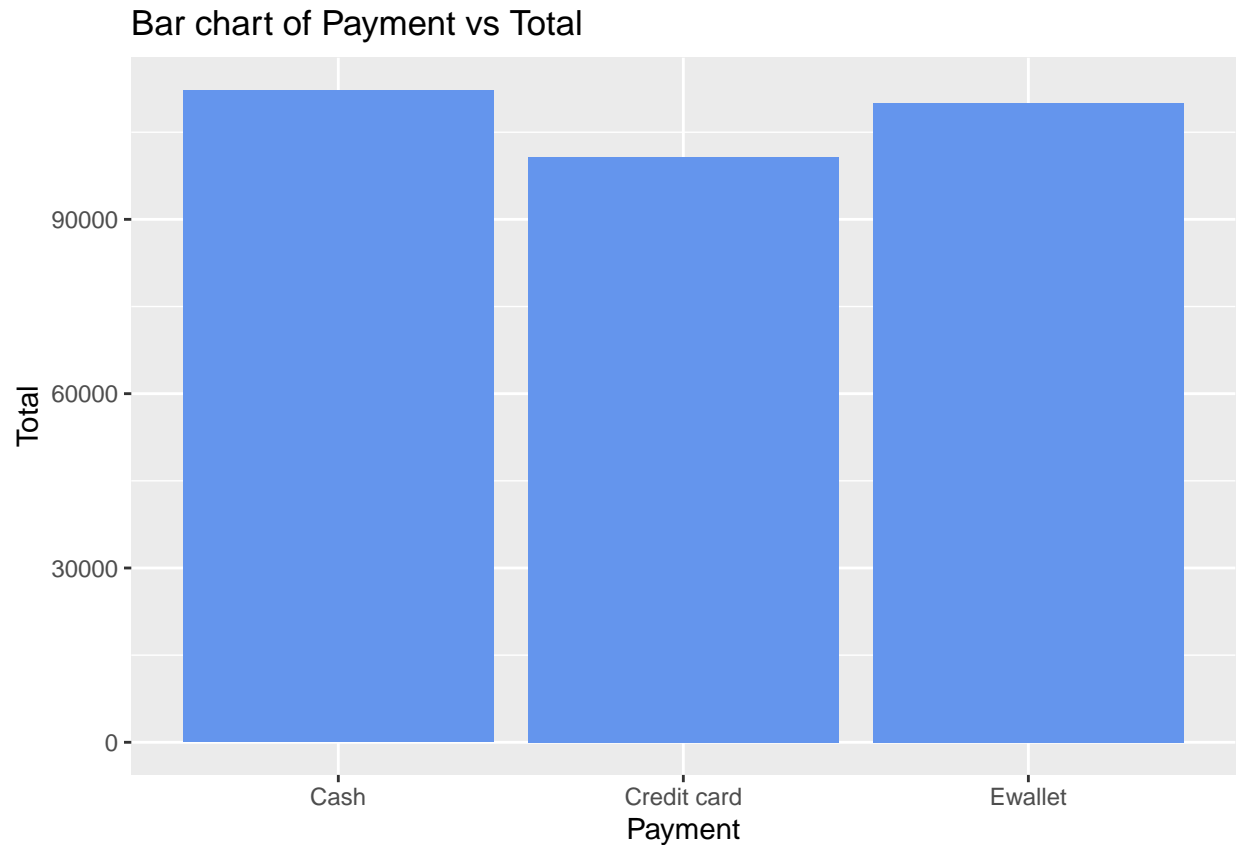


The branch with the highest number of sales is c

```
# Bar chart of Payment vs Total
agg_2 <- aggregate(smarket$Total,list(smarket$Payment),sum)
agg_2
```

```
##      Group.1      x
## 1      Cash 112206.6
## 2 Credit card 100767.1
## 3      Ewallet 109993.1
```

```
ggplot(data = agg_2 , aes(x= Group.1, y =x)) + geom_col(fill ="cornflowerblue") +
  labs(y="Total", x="Payment")+ggtitle("Bar chart of Payment vs Total")
```

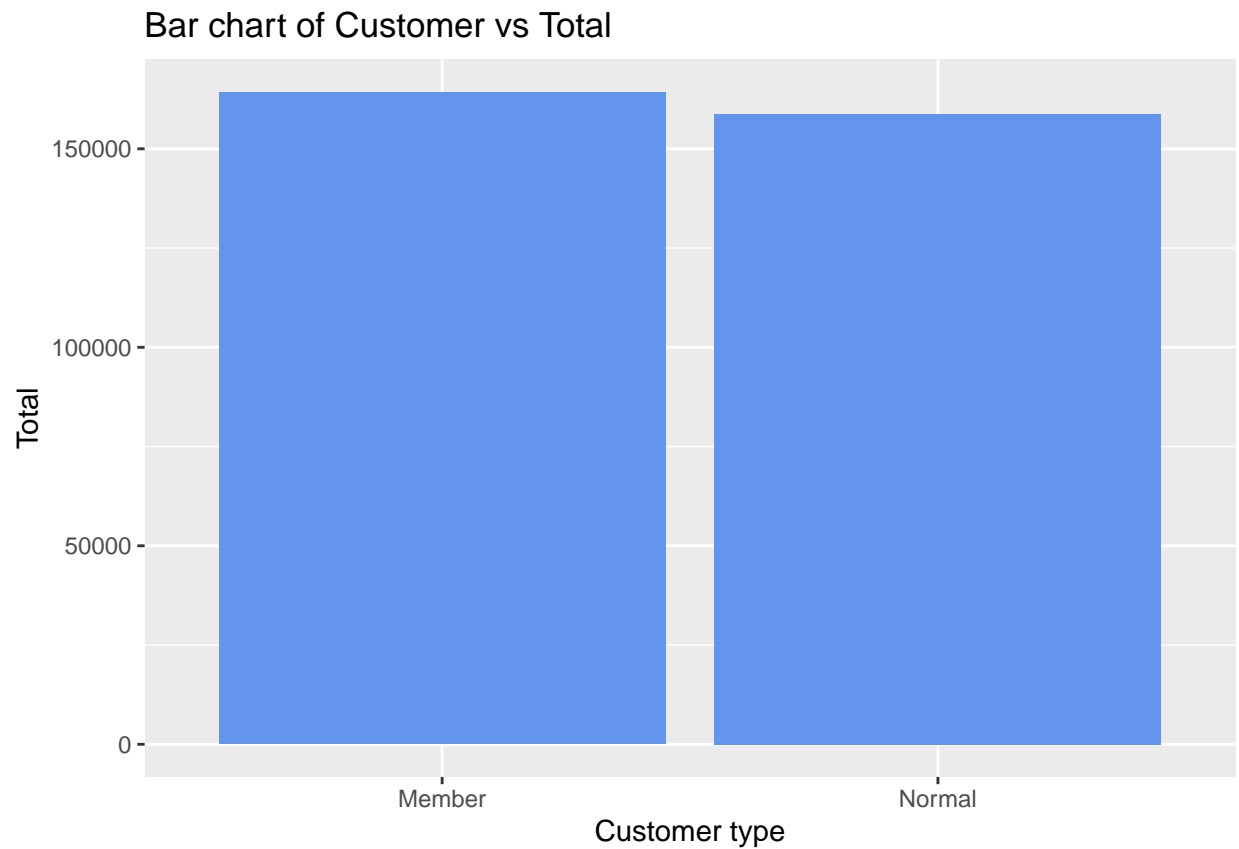


The Payment method with the highest number of sales is cash

```
# Bar chart of Customer type vs Total
agg_3 <- aggregate(smarket$Total,list(smarket$Customer.type),sum)
agg_3
```

```
##   Group.1      x
## 1  Member 164223.4
## 2  Normal 158743.3
```

```
ggplot(data = agg_3 , aes(x= Group.1, y =x)) + geom_col(fill ="cornflowerblue") +
  labs(y="Total", x="Customer type")+ggtitle("Bar chart of Customer vs Total")
```




Member customer type have the highest number of sales

Bar chart of Gender vs Total

```
agg_4 <- aggregate(smarket$Total,list(smarket$Gender),sum)
agg_4
```

```
##   Group.1      x
## 1  Female 167882.9
## 2   Male 155083.8
```

```
ggplot(data = agg_4 , aes(x= Group.1, y =x)) + geom_col(fill ="cornflowerblue") +
  labs(y="Total", x="Gender")+ggtitle("Bar chart of Gender vs Total")
```




Female customers have highest number of sales

```
# Bar chart of Product line vs Total
agg_5 <- aggregate(smarket$Total,list(smarket$Product.line),sum)
agg_5
```

```
##           Group.1      x
## 1 Electronic accessories 54337.53
## 2   Fashion accessories 54305.89
## 3   Food and beverages 56144.84
## 4   Health and beauty 49193.74
## 5   Home and lifestyle 53861.91
## 6   Sports and travel 55122.83
```

```
gg.agg <- ggplot (data = agg_5, aes(x= Group.1, y =x)) +
  geom_col(fill = "cornflowerblue") + theme(axis.text.x = element_text(
    angle = 90, vjust =.5, hjust = 1
  )) + labs(y="Total", x="Product line")+ggtitle("Bar chart of Product line vs Total")
gg.agg + ggtitle("Bar chart on Product Line")
```




Food and beverages have the highest number of sales

```
# Bar chart of Month vs Total
```

```
agg_6 <- aggregate(smarket$Total,list(smarket$Month),sum)
agg_6
```

```
##      Group.1      x
## 1 February  97219.37
## 2 January 116291.87
## 3   March 109455.51
```

```
ggplot(data = agg_6 , aes(x= Group.1, y =x)) + geom_col(fill ="cornflowerblue") +
  labs(y="Total", x="Month")+ggtitle("Bar chart of Month vs Total")
```



January is the month with the highest number of sales

Data Preprocessing

```
# We will start with converting our categorical data to numerical data type  
# Convert our variables to factor datatype  
final.market <- smarket  
final.market$Branch<- as.factor(final.market$Branch)  
final.market$Product.line <- as.factor(final.market$Product.line)  
final.market$Month <- as.factor(final.market$Month)  
final.market$Customer.type <- as.factor(final.market$Customer.type)  
final.market$Gender <- as.factor(final.market$Gender)  
final.market$Payment<- as.factor(final.market$Payment)
```

```
# Convert our variables from factor to numeric datatype  
final.market$Branch<- as.numeric(final.market$Branch)  
final.market$Product.line <- as.numeric(final.market$Product.line)  
final.market$Month <- as.numeric(final.market$Month)  
final.market$Customer.type <- as.numeric(final.market$Customer.type)  
final.market$Gender <- as.numeric(final.market$Gender)  
final.market$Payment<- as.numeric(final.market$Payment)
```

```
# Confirm changes made
head(final.market)
```

```
##      Invoice.ID Branch Customer.type Gender Product.line Unit.price Quantity
## 1 750-67-8428      1           1      1         4       74.69          7
## 2 226-31-3081      3           2      1         1       15.28          5
## 3 631-41-3108      1           2      2         5       46.33          7
## 4 123-19-1176      1           1      2         4       58.22          8
## 5 373-73-7910      1           2      2         6       86.31          7
## 6 699-14-3026      3           2      2         1       85.39          7
##      Tax      Date  Time Payment  cogs gross.margin.percentage gross.income
## 1 26.1415 2019-01-05 13:08      3 522.83          4.761905          26.1415
## 2  3.8200 2019-03-08 10:29      1  76.40          4.761905           3.8200
## 3 16.2155 2019-03-03 13:23      2 324.31          4.761905          16.2155
## 4 23.2880 2019-01-27 20:33      3 465.76          4.761905          23.2880
## 5 30.2085 2019-02-08 10:37      3 604.17          4.761905          30.2085
## 6 29.8865 2019-03-25 18:30      3 597.73          4.761905          29.8865
##      Rating      Total Month
## 1    9.1 548.9715      2
## 2    9.6  80.2200      3
## 3    7.4 340.5255      3
## 4    8.4 489.0480      2
## 5    5.3 634.3785      1
## 6    4.1 627.6165      3
```

```
# We will carry out our analysis without invoice variable because it's a
# unique variables, gross.margin.percentage variable only has one value hence
# 0 variance, instead of using the date and time variable we will use month
final <- final.market[,c(-1,-9,-10,-13)]
# View the new dataset
head(final)
```

```
##      Branch Customer.type Gender Product.line Unit.price Quantity      Tax Payment
## 1      1           1      1         4       74.69          7 26.1415          3
## 2      3           2      1         1       15.28          5  3.8200          1
## 3      1           2      2         5       46.33          7 16.2155          2
## 4      1           1      2         4       58.22          8 23.2880          3
## 5      1           2      2         6       86.31          7 30.2085          3
## 6      3           2      2         1       85.39          7 29.8865          3
##      cogs gross.income Rating      Total Month
## 1 522.83      26.1415    9.1 548.9715      2
## 2  76.40       3.8200    9.6  80.2200      3
## 3 324.31      16.2155    7.4 340.5255      3
## 4 465.76      23.2880    8.4 489.0480      2
## 5 604.17      30.2085    5.3 634.3785      1
## 6 597.73      29.8865    4.1 627.6165      3
```

Feature Selection


```
# Loading the relevant packages
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.1.3
```

```
## Loading required package: lattice
```

```
##
```

```
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 4.1.3
```

```
## corrplot 0.92 loaded
```

Filter Method

```
# Calculating the correlation matrix
corr_final <- cor(final[, -12])
corr_final
```

```
##           Branch Customer.type      Gender Product.line  Unit.price
## Branch      1.00000000 -0.01960787 -0.056317558 -0.053937557 0.028202440
## Customer.type -0.01960787  1.00000000  0.039996160 -0.036800311 -0.020237875
## Gender      -0.05631756  0.03999616  1.000000000  0.005193197 0.015444630
## Product.line -0.05393756 -0.03680031  0.005193197  1.000000000 0.019321028
## Unit.price   0.02820244 -0.02023787  0.015444630  0.019321028 1.000000000
## Quantity     0.01596379 -0.01676271 -0.074258307  0.020256001 0.010777564
## Tax          0.04104666 -0.01967028 -0.049450989  0.031620725 0.633962089
## Payment     -0.05010429  0.01807344  0.044577609  0.029896383 -0.015941048
## cogs         0.04104666 -0.01967028 -0.049450989  0.031620725 0.633962089
## gross.income 0.04104666 -0.01967028 -0.049450989  0.031620725 0.633962089
## Rating       0.01023848  0.01888867  0.004800208 -0.020528973 -0.008777507
## Month        -0.04033140  0.04488694  0.059814418  0.044994245 -0.034715674
##           Quantity      Tax      Payment      cogs gross.income
## Branch      0.015963788 0.041046665 -0.050104288 0.041046665 0.041046665
## Customer.type -0.016762706 -0.019670283 0.018073436 -0.019670283 -0.019670283
## Gender      -0.074258307 -0.049450989 0.044577609 -0.049450989 -0.049450989
## Product.line 0.020256001 0.031620725 0.029896383 0.031620725 0.031620725
## Unit.price   0.010777564 0.633962089 -0.015941048 0.633962089 0.633962089
## Quantity     1.000000000 0.705510186 -0.003920990 0.705510186 0.705510186
## Tax          0.705510186 1.000000000 -0.012433637 1.000000000 1.000000000
## Payment     -0.003920990 -0.012433637 1.000000000 -0.012433637 -0.012433637
```

```
## cogs          0.705510186  1.000000000 -0.012433637  1.000000000  1.000000000
## gross.income  0.705510186  1.000000000 -0.012433637  1.000000000  1.000000000
## Rating        -0.015814905 -0.036441705 -0.005381289 -0.036441705 -0.036441705
## Month         0.002375544 -0.006724205  0.041898766 -0.006724205 -0.006724205
##              Rating      Month
## Branch         0.010238476 -0.040331403
## Customer.type  0.018888672  0.044886937
## Gender         0.004800208  0.059814418
## Product.line   -0.020528973  0.044994245
## Unit.price     -0.008777507 -0.034715674
## Quantity       -0.015814905  0.002375544
## Tax            -0.036441705 -0.006724205
## Payment        -0.005381289  0.041898766
## cogs           -0.036441705 -0.006724205
## gross.income   -0.036441705 -0.006724205
## Rating         1.000000000 -0.054714932
## Month         -0.054714932  1.000000000
```

```
# Find attributes that are highly correlated
# Cut off value of 0.70
high.corr <- findCorrelation(corr_final, cutoff=0.70)
high.corr
```

```
## [1] 7 9 10
```

```
# Highly correlated attributes
names(final[,high.corr])
```


```
## [1] "Tax"          "cogs"          "gross.income"
```

Variables with higher correlation are Tax , Cogs and gross income

```
# We will remove the variables with a higher correlation
# and compare the results graphically

# Removing Highly correlated attributes Features
final_1 <- final[-high.corr]

# Performing our graphical comparison
par(mfrow = c(1, 2))
corrplot(corr_final, order = "hclust")
corrplot(cor(final_1), order = "hclust")
```



Wrapper Method

```
# Loading packages
library(clustvarsel)

## Warning: package 'clustvarsel' was built under R version 4.1.3

## Loading required package: mclust

## Warning: package 'mclust' was built under R version 4.1.3

## Package 'mclust' version 5.4.9
## Type 'citation("mclust")' for citing this R package in publications.

##
## Attaching package: 'mclust'

## The following object is masked from 'package:purrr':
##
##   map

## Package 'clustvarsel' version 2.3.4
```

```
## Type 'citation("clustvarsel")' for citing this R package in publications.
```

```
library(mclust)
# Sequential forward greedy search (default)
frwr = clustvarsel(final[, -12], G = 1:5)
frwr
```

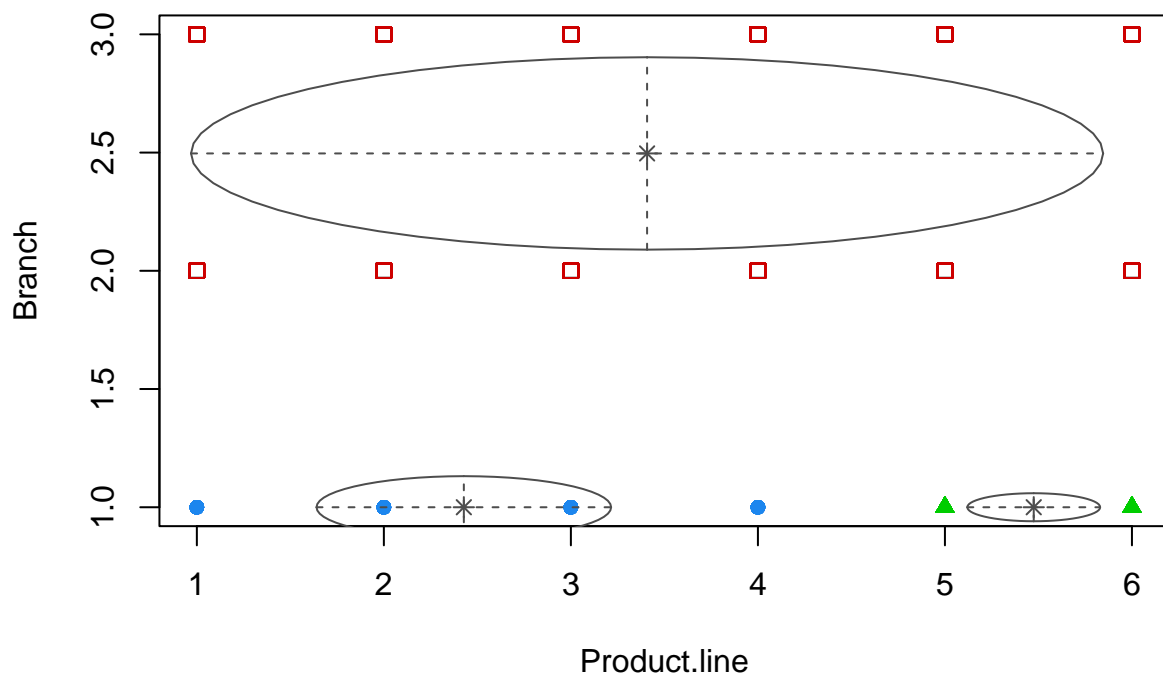
```
## -----
## Variable selection for Gaussian model-based clustering
## Stepwise (forward/backward) greedy search
## -----
##
## Variable proposed Type of step BICclust Model G BICdiff Decision
## Product.line Add -3521.631 E 5 408.36743 Accepted
## Branch Add -5233.923 VEV 5 735.71798 Accepted
## Customer.type Add -6722.774 VEV 4 -23.45627 Rejected
## Branch Remove -3498.098 E 5 712.18492 Rejected
##
## Selected subset: Product.line, Branch
```

The selection algorithm indicates that the subset we use for the clustering model is composed of variables product line and Branch

```
# Let's build clustering model:
Subset = final[, -12][, frwr$subset]
mod = Mclust(Subset, G = 1:5)
summary(mod)
```

```
## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
##
## Mclust VEI (diagonal, equal shape) model with 3 components:
##
## log-likelihood n df BIC ICL
## -2723.585 1000 12 -5530.063 -5531.496
##
## Clustering table:
## 1 2 3
## 216 660 124
```

```
# Plot the clustering model
plot(mod, c("classification"))
```



Feature Ranking

```
# Import the library
library(FSelector)
```

```
## Warning: package 'FSelector' was built under R version 4.1.3
```

```
# We'll use the correlation coefficient as a unit of valuation.
fr <- linear.correlation(Total~., final)
fr
```

```
##          attr_importance
## Branch          0.041046665
## Customer.type    0.019670283
## Gender           0.049450989
## Product.line     0.031620725
## Unit.price       0.633962089
## Quantity         0.705510186
## Tax              1.000000000
## Payment          0.012433637
## cogs             1.000000000
## gross.income     1.000000000
```

```
## Rating          0.036441705
## Month           0.006724205
```

```
# Top five representative variables
var <- cutoff.k(fr, 5)
as.data.frame(var)
```

```
##          var
## 1         Tax
## 2         cogs
## 3 gross.income
## 4      Quantity
## 5    Unit.price
```

```
# Instead of using the scores for the correlation coefficient,
# we can use information gain
fr.1 <- information.gain(Total~., final)
```

```
# Top five representative variables
var.1 <- cutoff.k(fr.1, 5)
as.data.frame(var.1)
```

```
##          var.1
## 1         Tax
## 2         cogs
## 3 gross.income
## 4      Quantity
## 5    Unit.price
```

The top 5 representative variables are the same for both correlation coefficient and information gain scores. They are Tax,cogs,gross.income,quantity and unit.price.

Embedded Methods

```
# load the wskm package
library(wskm)
```

```
## Warning: package 'wskm' was built under R version 4.1.3
```

```
## Loading required package: latticeExtra
```

```
## Warning: package 'latticeExtra' was built under R version 4.1.3
```

```
##
## Attaching package: 'latticeExtra'
```

```
## The following object is masked from 'package:ggplot2':
##
##      layer
```

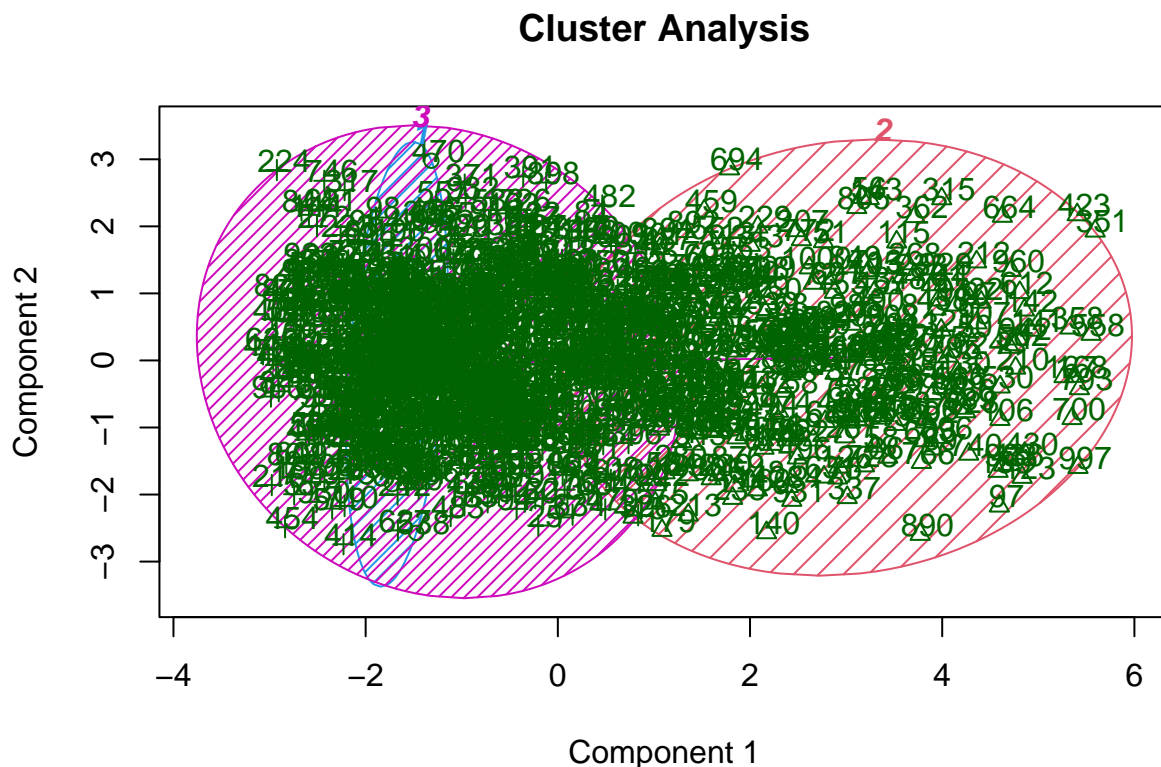
```
## Loading required package: fpc
```

```
## Warning: package 'fpc' was built under R version 4.1.3
```

```
# Set up the model
set.seed(2)
model <- ewkm(final[, -12], 3, lambda=2, maxiter=1000)
```

```
# Loading and installing our cluster package
library("cluster")
```

```
# Cluster Plot against 1st 2 principal components
clusplot(final[, -12], model$cluster, color=TRUE, shade=TRUE,
          labels=2, lines=1, main='Cluster Analysis')
```



These two components explain 42.69 % of the point variability.

```
# Weights are calculated for each variable and cluster.
# They are a measure of the relative importance of each variable
# with regards to the membership of the observations to that cluster.
# The weights are incorporated into the distance function,
# typically reducing the distance for more important variables.
# Weights remain stored in the model and we can check them as follows:
round(model$weights*100,2)
```

```
## Branch Customer.type Gender Product.line Unit.price Quantity Tax Payment
```

## 1	0	0.13	0.14	0	0	0	49.86	0
## 2	0	48.12	51.87	0	0	0	0.00	0
## 3	0	50.00	50.00	0	0	0	0.00	0
##	cogs	gross.income	Rating	Month				
## 1	0	49.86	0	0				
## 2	0	0.00	0	0				
## 3	0	0.00	0	0				

Customer type and gender hold weight across the 3 clusters.